

CMPT412 Project 3

Digit recognition with convolutional neural networks

Student: Sibei Zhou(301416917)

Computer ID: palamzad@sfsu.ca

Team members: Sibei Zhou, Honghao Yu, ChenKun Zhang

Kaggle entry name: Sibei Zhou

(use 3 free-late days)

- Part 1:

- List of the configs and modifications that you used.

- (1) Default configs:

```
cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url  
    ("COCO-Detection/faster_rcnn_R_101_FPN_3x.yaml")  
  
cfg.SOLVERIMS_PER_BATCH = 2  
  
cfg.SOLVER.BASE_LR = 0.00025 #0.00025->0.0005->0.001->0.0005(final)  
  
cfg.SOLVER.MAX_ITER = 500 # 500->1000->5000(final)  
  
cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 512  
  
cfg.MODEL.ROI_HEADS.NUM_CLASSES = 1
```

- (2) Final configs: (Modified the MAX_ITER, BASE_LR and change the default model to faster_rcnn_X_101_32x8d_FPN_3x.yaml)

```
cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url  
    ("COCO-Detection/faster_rcnn_X_101_32x8d_FPN_3x.yaml")  
  
cfg.SOLVERIMS_PER_BATCH = 2  
  
cfg.SOLVER.BASE_LR = 0.0005  
  
cfg.SOLVER.MAX_ITER = 5000  
  
cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 512  
  
cfg.MODEL.ROI_HEADS.NUM_CLASSES = 1
```

- Factors which helped improve the performance. Explain each factor in 2-3 lines.

■MAX_ITER

Since the loss is still decreasing rapidly when max_iter = 500, the higher iterations might improve the model and get higher accuracy, so we keep trying on more iteration until total_loss becomes stable.

■BASE_LR

We have modified the learning rate a bit higher (0.0005) and higher (0.001) than the original (0.00025), the accuracy of lr=0.0005 has increased but another with lr=0.001 got a very low accuracy. So we choose a learning rate 0.0005 at the end.

Result compared with different learning rate at MAX_ITER = 500:

BASE_LR = 0.00025 (default) :

```
[03/01 04:04:23 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | AP1 |
|-----|-----|-----|-----|-----|-----|
| 29.045 | 50.717 | 30.527 | 20.255 | 36.919 | 59.305 |
```

BASE_LR = 0.0005 (best) :

```
[03/01 05:34:18 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | AP1 |
|-----|-----|-----|-----|-----|-----|
| 32.265 | 55.007 | 35.175 | 22.874 | 40.727 | 59.979 |
```

BASE_LR = 0.001 (worse) :

```
[03/01 05:00:23 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | AP1 |
|-----|-----|-----|-----|-----|-----|
| 5.088 | 13.425 | 2.088 | 6.700 | 10.433 | 11.504 |
```

■MODEL.WEIGHTS

On the Model_ZOO.md of Detectron2 we found in github, it introduces the X101-FPN model has a slightly higher box AP compared to the original R101-FPN model (picture below), so we modified the model to use the X101-FPN.

R101-FPN	3x	0.286	0.051	4.1	42.0	137851257	model metrics
X101-FPN	3x	0.638	0.098	6.7	43.0	139173657	model metrics

- Final plot for total training loss and accuracy. This would have been auto-generated by the notebook.

```
[03/01 10:49:08 d2.utils.events]: eta: 0:03:57 iter: 4519 total_loss: 0.4471
[03/01 10:49:19 d2.utils.events]: eta: 0:03:47 iter: 4539 total_loss: 0.4143
[03/01 10:49:30 d2.utils.events]: eta: 0:03:36 iter: 4559 total_loss: 0.4004
[03/01 10:49:41 d2.utils.events]: eta: 0:03:26 iter: 4579 total_loss: 0.4839
[03/01 10:49:53 d2.utils.events]: eta: 0:03:16 iter: 4599 total_loss: 0.3368
[03/01 10:50:04 d2.utils.events]: eta: 0:03:06 iter: 4619 total_loss: 0.3388
[03/01 10:50:16 d2.utils.events]: eta: 0:02:56 iter: 4639 total_loss: 0.4339
[03/01 10:50:28 d2.utils.events]: eta: 0:02:47 iter: 4659 total_loss: 0.3659
[03/01 10:50:40 d2.utils.events]: eta: 0:02:37 iter: 4679 total_loss: 0.4107
[03/01 10:50:51 d2.utils.events]: eta: 0:02:27 iter: 4699 total_loss: 0.3912
[03/01 10:51:02 d2.utils.events]: eta: 0:02:18 iter: 4719 total_loss: 0.4032
[03/01 10:51:13 d2.utils.events]: eta: 0:02:08 iter: 4739 total_loss: 0.3949
[03/01 10:51:24 d2.utils.events]: eta: 0:01:58 iter: 4759 total_loss: 0.4526
[03/01 10:51:36 d2.utils.events]: eta: 0:01:48 iter: 4779 total_loss: 0.3394
[03/01 10:51:46 d2.utils.events]: eta: 0:01:38 iter: 4799 total_loss: 0.4053
[03/01 10:51:57 d2.utils.events]: eta: 0:01:28 iter: 4819 total_loss: 0.4539
[03/01 10:52:09 d2.utils.events]: eta: 0:01:19 iter: 4839 total_loss: 0.4308
[03/01 10:52:20 d2.utils.events]: eta: 0:01:09 iter: 4859 total_loss: 0.3203
[03/01 10:52:32 d2.utils.events]: eta: 0:00:59 iter: 4879 total_loss: 0.3234
[03/01 10:52:43 d2.utils.events]: eta: 0:00:49 iter: 4899 total_loss: 0.3849
[03/01 10:52:55 d2.utils.events]: eta: 0:00:39 iter: 4919 total_loss: 0.3821
[03/01 10:53:06 d2.utils.events]: eta: 0:00:29 iter: 4939 total_loss: 0.2878
[03/01 10:53:16 d2.utils.events]: eta: 0:00:19 iter: 4959 total_loss: 0.3816
[03/01 10:53:27 d2.utils.events]: eta: 0:00:09 iter: 4979 total_loss: 0.3736
[03/01 10:53:41 d2.utils.events]: eta: 0:00:00 iter: 4999 total_loss: 0.379
```

Average Precision	(AP) @ [IoU=0.50:0.95]	area=	all	maxDets=100] = 0.486	
Average Precision	(AP) @ [IoU=0.50]	area=	all	maxDets=100] = 0.695	
Average Precision	(AP) @ [IoU=0.75]	area=	all	maxDets=100] = 0.561	
Average Precision	(AP) @ [IoU=0.50:0.95]	area=	small	maxDets=100] = 0.387	
Average Precision	(AP) @ [IoU=0.50:0.95]	area=	medium	maxDets=100] = 0.553	
Average Precision	(AP) @ [IoU=0.50:0.95]	area=	large	maxDets=100] = 0.821	
Average Recall	(AR) @ [IoU=0.50:0.95]	area=	all	maxDets= 1] = 0.020	
Average Recall	(AR) @ [IoU=0.50:0.95]	area=	all	maxDets= 10] = 0.164	
Average Recall	(AR) @ [IoU=0.50:0.95]	area=	all	maxDets=100] = 0.511	
Average Recall	(AR) @ [IoU=0.50:0.95]	area=	small	maxDets=100] = 0.392	
Average Recall	(AR) @ [IoU=0.50:0.95]	area=	medium	maxDets=100] = 0.587	
Average Recall	(AR) @ [IoU=0.50:0.95]	area=	large	maxDets=100] = 0.870	
[03/01 10:55:28 d2.evaluation.coco_evaluation]: Evaluation results for bbox:					
AP	AP50	AP75	APs	APm	AP1
48.584	69.491	56.087	38.702	55.279	82.125

- The visualization of 3 samples from the test set and the predicted results.
 - Sample test 1



- Sample test 2



- Sample test 3



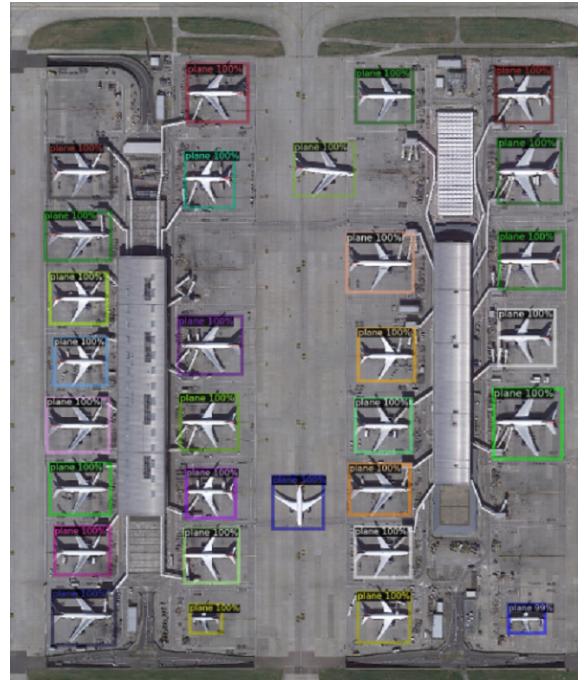
- At least one ablation study to validate the above choices, i.e., a comparison of performance for two variants of a model, one with and one without a certain feature or implementation choice. In addition, provide visualization of a sample from the test set for qualitative comparison.
 - Model with `MAX_ITER = 1000`

With the lower iterations, the training process only takes 10 mins, and the result looks great with only a few plane missed (left figure):



- Model with `MAX_ITER = 5000`

When the iterations increase, the training process takes over 30 mins, but it got lower `total_loss` and the result looks greater than before (right figure):



- Part 2:

- Report any hyperparameter settings you used (batch_size, learning_rate, num_epochs, optimizer).

```

num_epochs = 50

batch_size = 8

learning_rate = 0.05      # 0.01 -> 0.05

weight_decay = 1e-5

optim = torch.optim.Adam(model.parameters(), lr=learning_rate,
weight_decay=weight_decay)

```

- Report the final architecture of your network including any modification that you have for the layers. Briefly explain the reason for each modification.

```

# Encoder

self.input_conv = conv(3, 8)
self.down_1 = down(8, 16)
self.down_2 = down(16, 32)
self.down_3 = down(32, 64)
self.down_4 = down(64, 128)
self.down_5 = down(128, 256)
self.down_6 = down(256, 512)

# Decoder

self.up_1 = up(512, 256)
self.up_2 = up(256, 128)
self.up_3 = up(128, 64)
self.up_4 = up(64, 32)
self.up_5 = up(32, 16)
self.up_6 = up(16, 8)
self.up_7 = conv(8, 3)
self.output_conv = conv(3, 1, False)

self.normal_1 = nn.BatchNorm2d(8)
self.normal_2 = nn.BatchNorm2d(16)
self.normal_3 = nn.BatchNorm2d(32)
self.normal_4 = nn.BatchNorm2d(64)
self.normal_5 = nn.BatchNorm2d(128)
self.normal_6 = nn.BatchNorm2d(256)

```

```

def forward(self, input):
    y_1 = self.input_conv(input)
    y_2 = self.down_1(y_1)
    y_3 = self.down_2(y_2)
    y_4 = self.down_3(y_3)
    y_5 = self.down_4(y_4)
    y_6 = self.down_5(y_5)
    y_7 = self.down_6(y_6)
    y_7 = self.normal_6(F.relu(self.up_1(y_7) + y_6))
    y_6 = self.normal_5(F.relu(self.up_2(y_7) + y_5))
    y_5 = self.normal_4(F.relu(self.up_3(y_6) + y_4))
    y_4 = self.normal_3(F.relu(self.up_4(y_5) + y_3))
    y_3 = self.normal_2(F.relu(self.up_5(y_4) + y_2))
    y_2 = self.normal_1(F.relu(self.up_6(y_3) + y_1))
    y = self.up_7(y_2)
    output = self.output_conv(y)
    return output

```

The model's architecture we used references the RNN from the last project, and we also add some skip layers that references the resnet on each decoder layer, which helps in improving accuracy a lot.

- Report the loss functions that you used and the plot the total training loss of the training procedure

```
crit = nn.BCEWithLogitsLoss() # Define the loss function
```

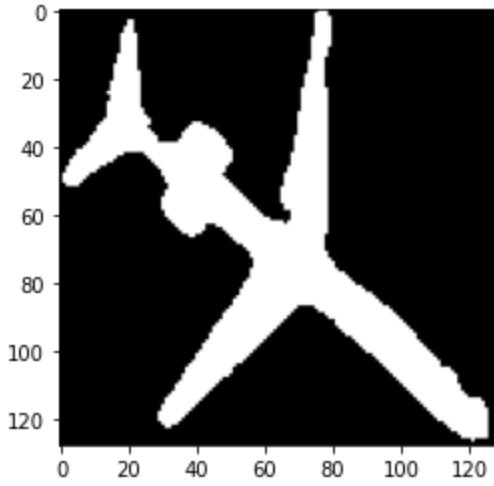
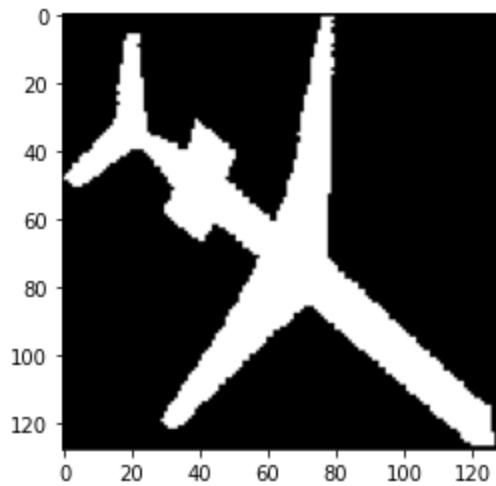
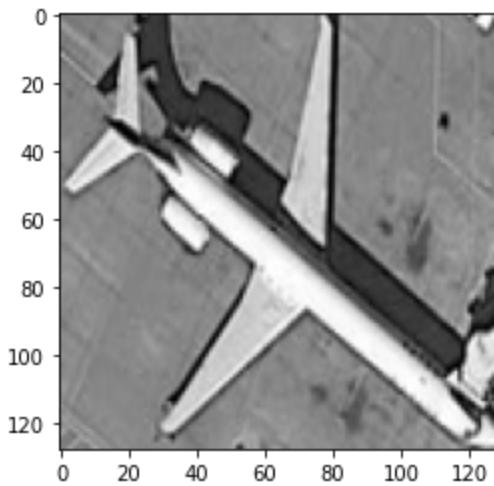
Loss: 0.050518546253442764

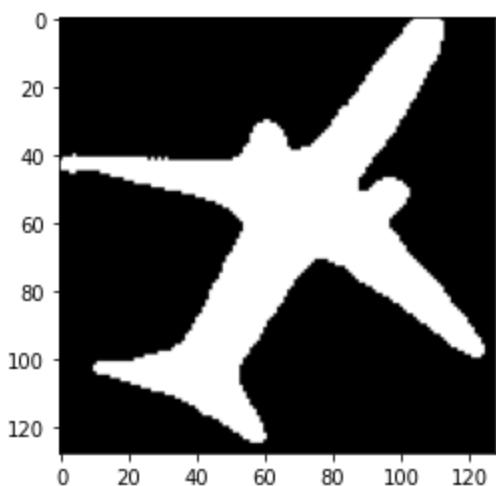
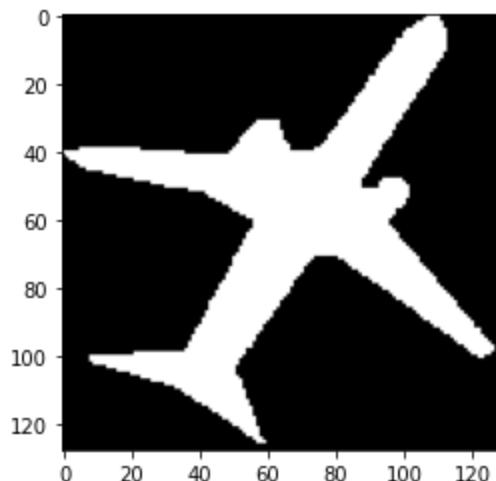
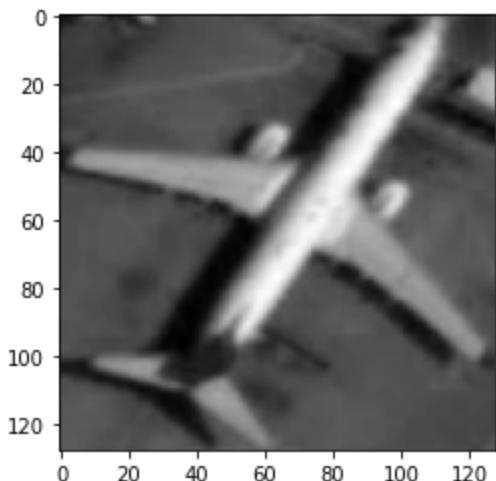
- Report the final mean IoU of your model

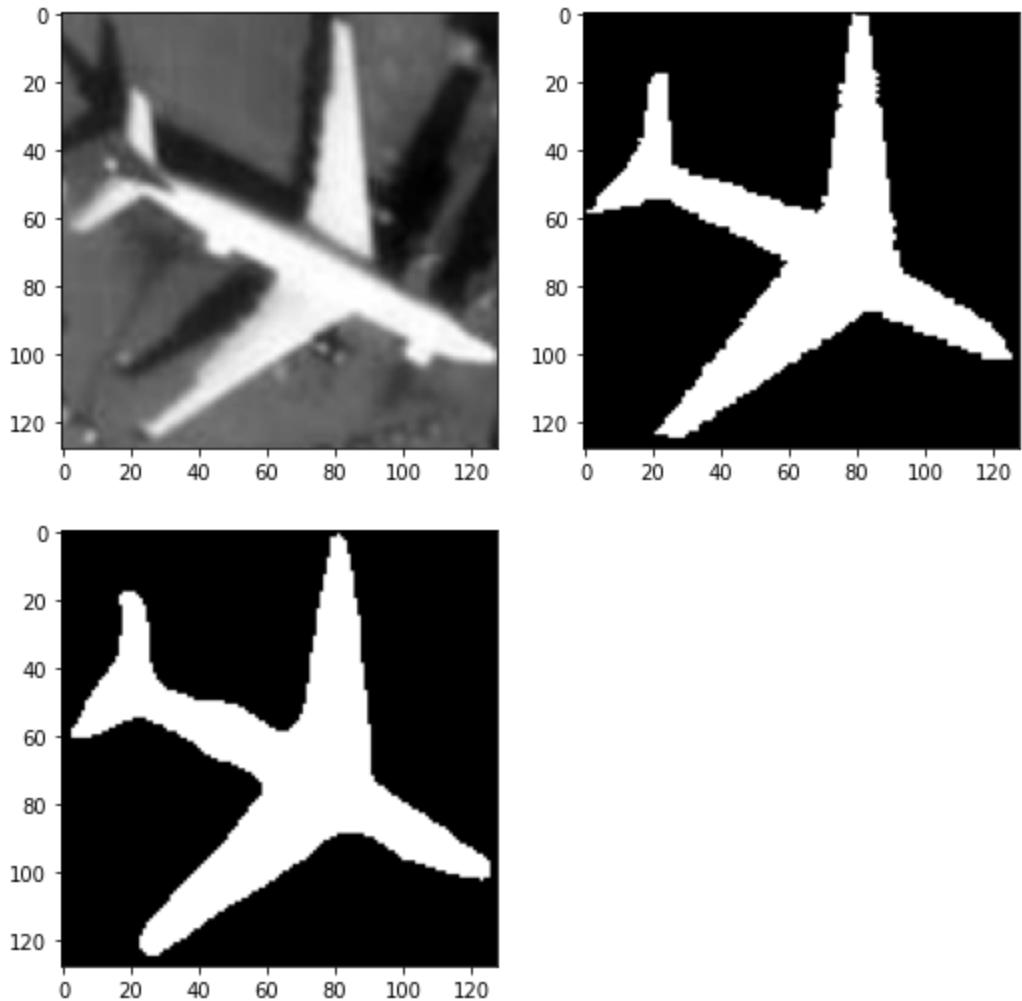
Mean IoU = 0.8718697733076178

- Visualize 3 images from the test set and the corresponding predicted masks

(on the right of the original picture is the mask picture, and the one below it is the prediction mask)







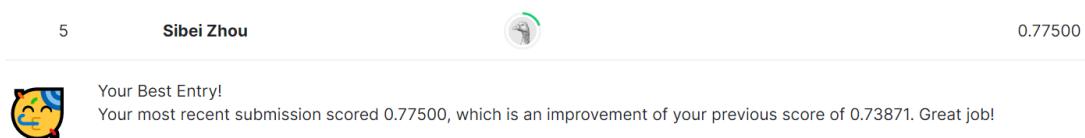
- Part 3:
 - The name under which you submitted on Kaggle. Names of your group members if you form a group.

Name on Kaggle: Sibei Zhou

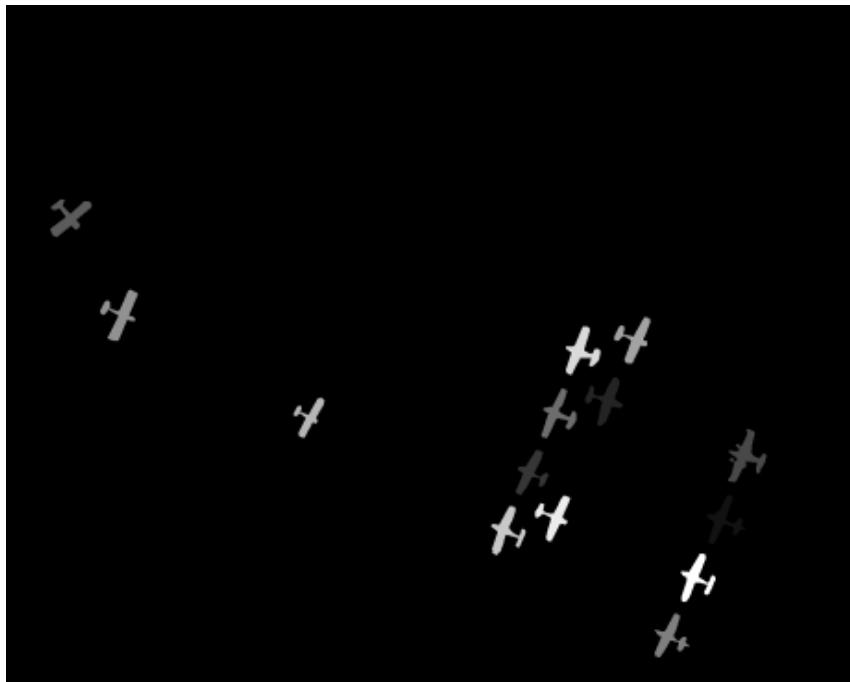
Groups Members: Sibei Zhou, Honghao Yu, Chenkun Zhang

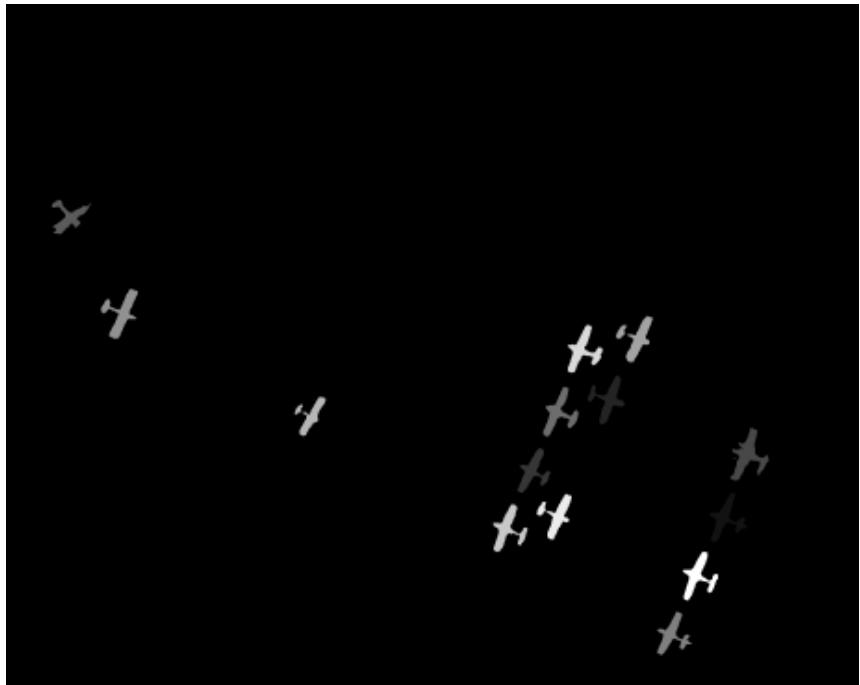
- Report the best score (should match your score on Kaggle).

Best score on kaggle: 0.775 (# 5 in Leaderboard)



- The visualization of results for 3 random samples from the test set.
(The first: original, the second: true mask, the third: predict mask)
 - Sample 1:





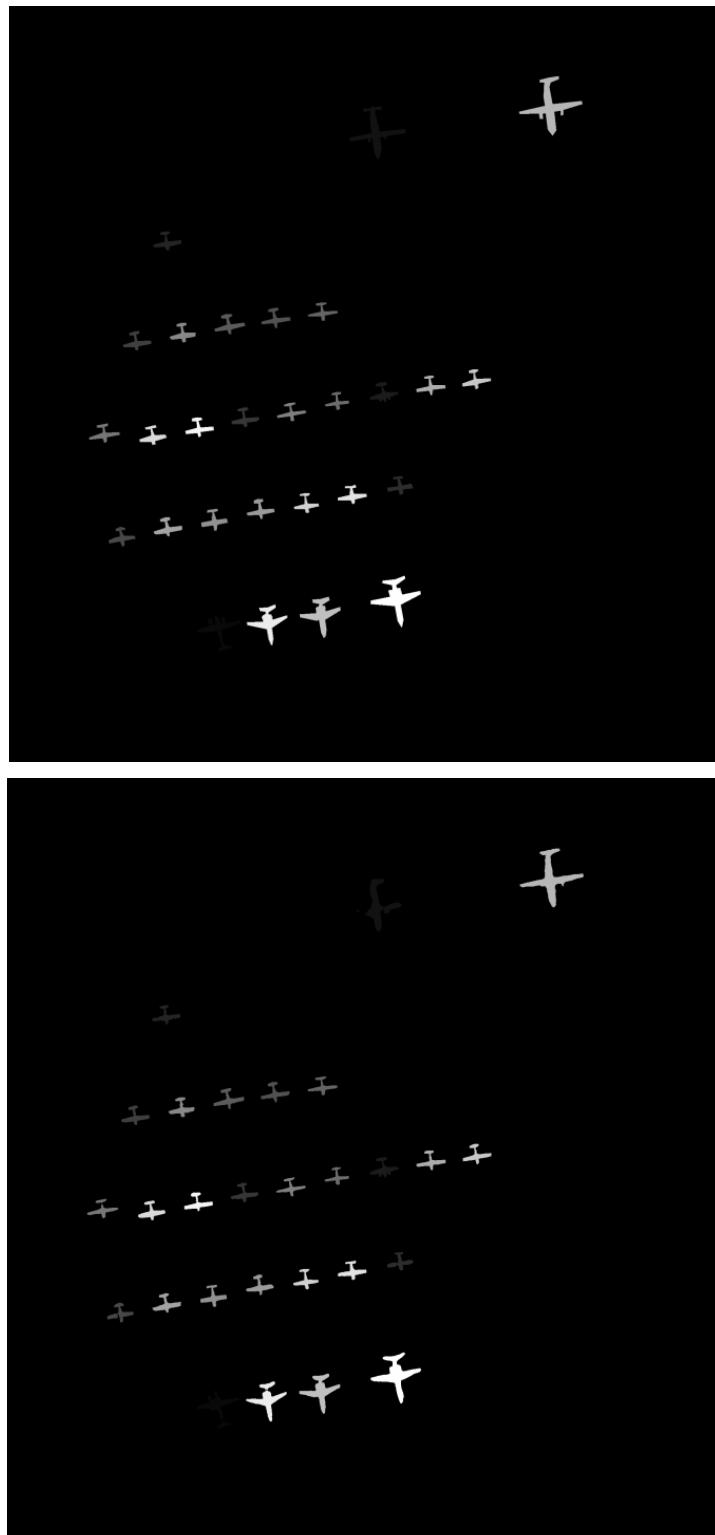
■ Sample 2:





■ Sample 3:





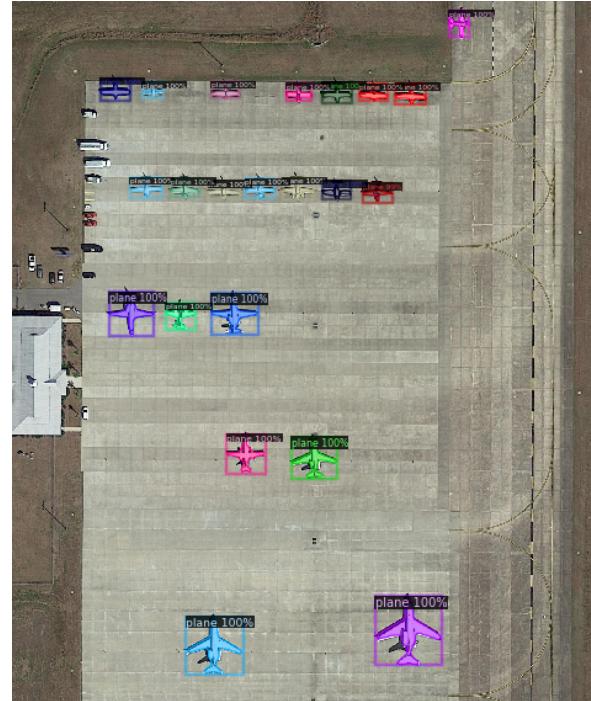
- CSV file of your predicted test labels needs to be uploaded to Kaggle.

- Part 4:

- The visualization and the evaluation results similar to Part 1
The hyperparameter used is the same as in Part 1, which are:

```
cfg.SOLVERIMS_PER_BATCH = 2  
cfg.SOLVER.BASE_LR = 0.0005  
cfg.SOLVER.MAX_ITER = 5000  
cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 512  
cfg.MODEL.ROI_HEADS.NUM_CLASSES = 1
```

visualization of 3 random samples:





- Explain the differences between the results of Part 3 and Part 4 in a few lines.

The result of Part 4 got lower accuracy than the result in part3 in both bbox AP:

Part 3 bbox AP:

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.486
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.695
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.561
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.387
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.553
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.821
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.020
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.164
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.511
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.392
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.587
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.870
[03/01 10:55:28 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | AP1 |
|-----|-----|-----|-----|-----|-----|
| 48.584 | 69.491 | 56.087 | 38.702 | 55.279 | 82.125 |

```

Part 4 bbox AP: 44.929

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.449
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.672
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.517
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.355
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.520
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.758
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.018
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.155
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.479
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.362
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.557
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.818
[03/02 11:21:43 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | AP1 |
|-----|-----|-----|-----|-----|-----|
| 44.929 | 67.197 | 51.741 | 35.464 | 51.965 | 75.786 |
```

And its performance in high-precision pictures is also worse than Part 3, in some high-precision pictures, it segmentation looks bad, for example:

