

## Project 2 Report

Sibei Zhou(301416917)

Submission name on Kaggle: Sibei Zhou

Team name: WOW

Best accuracy: **64.6%**

### Part 1


Layer Structure Table (The color of the table is to better distinguish the groups of the layers, like the convolutional modules(conv-relu-maxpool))

Layer No.	Layer Type	Kernel size	Input Output Dimension	Input Output Channels
1	Conv2d	3	32 32	3 64
2	Batchnorm	-	32 32	-
3	Relu	-	32 32	-
4	Conv2d	3	32 32	64 64
5	Batchnorm	-	32 32	-
6	Relu	-	32 32	-
7	Conv2d	3	32 32	64 128
8	Batchnorm	-	32 32	-
9	Relu	-	32 32	-


10	Conv2d	3	32 32	128 128
11	Batchnorm	-	32 32	-
12	Relu	-	32 32	-
13	Maxpool2d	2	32 16	-
14	Conv2d	3	16 16	128 192
15	Batchnorm	-	16 16	-
16	Relu	-	16 16	-
17	Conv2d	3	16 16	192 384
18	Batchnorm	-	16 16	-
19	Relu	-	16 16	-
20	Maxpool2d	2	16 8	-
21	Conv2d	3	8 8	384 256
22	Batchnorm	-	8 8	-
23	Relu	-	8 8	-
24	Conv2d	3	8 8	256 256
25	Batchnorm	-	8 8	-
26	Relu	-	8 8	-

27	Maxpool2d	2	8 4	-
28	Linear	-	4096 2048	-
29	Relu	-	2048 2048	-
30	Linear	-	2048 1024	-
31	Relu	-	1024 1024	-
32	Dropout(p=0.2)		1024 1024	
33	Linear	-	1024 512	-
34	Relu	-	512 512	-
35	Linear	-	512 216	-
36	Relu	-	216 216	-
37	Linear	-	216 100	-

**Result plot (epoch = 20)**



## Result plot (epoch = 40)



We can see in the graph that even though the train loss keeps decreasing, the val accuracy only slightly improved after rising to 60%, and then stabilized on the last 10 epochs.

## Structure Explain

Part of the structure is referenced from the AlexNet model, but many modifications have been made to make the net tailored for picture size 32\*32.

First, the number of convolution layers is increased to 8 from original AlexNet's 5 convolution layers, with increasing output channels, which prevents the net from overfitting and reduces the training accuracy.

Then, a 2d batch normalizing layer is added before the RELU layer of each convolution layer, and Pooling layers with 2 kernel size and stride = 1 are used instead, and we put it after the 4th, 6th, 8th convolution layer.

Moreover, one more fully connected layer after AlexNet.

## Ablation Study

1. Adding 3 more convolution layers with 32,64,128,256 output channels and change all convolution layers' kernel size to 3: accuracy increase **from 26.3% to 35.8%**
2. Batching normalizing layer added after each convolution layer, and Relu layers after it: accuracy increase **from 35.8% to 38%**

3. Adapting AlexNet structure, with 5 convolution layers + Relu layers, 3 pooling layers(after the 1st, 2nd, 5th convolution layer), and 2 linear layers + Relu + Dropout(p=0.5), and the final linear layer which return the result label between 100 (in this step, since the original AlexNet use convolution layer with kernel size = 11, it is too big and reduce the output dimension too fast, and lead to failure, so we change all the size of the kernel to 3): accuracy increase **from 35.8% to 43.4%**
4. Change the portion of the 3 pooling layers and change the convolution layers' filter size: accuracy increase **from 43.4% to 62.4%**
5. Increase the epoch from 20 to 60: accuracy increase **from 62.4% to 64.6%**

## Part 2

### Model and Hyperparameters

The first time I just use the pre-training renet18 without changing the given hyperparameters, I achieve a train accuracy of 14.9%:

```
TRAINING Epoch 1/10 Loss 0.6792 Accuracy 0.0043
TRAINING Epoch 2/10 Loss 0.6617 Accuracy 0.0110
TRAINING Epoch 3/10 Loss 0.6498 Accuracy 0.0180
TRAINING Epoch 4/10 Loss 0.6382 Accuracy 0.0273
TRAINING Epoch 5/10 Loss 0.6279 Accuracy 0.0460
TRAINING Epoch 6/10 Loss 0.6181 Accuracy 0.0547
TRAINING Epoch 7/10 Loss 0.6086 Accuracy 0.0780
TRAINING Epoch 8/10 Loss 0.5989 Accuracy 0.0953
TRAINING Epoch 9/10 Loss 0.5893 Accuracy 0.1220
TRAINING Epoch 10/10 Loss 0.5797 Accuracy 0.1493
Finished Training
```

1. First try on slight tweaks to those hyperparameters, that is:

```
NUM_EPOCHS = 50

LEARNING_RATE = 0.001

BATCH_SIZE = 32

RESNET_LAST_ONLY = True
```

I get a train accuracy 53.67%.

```
TRAINING Epoch 40/50 Loss 0.0846 Accuracy 0.5157
TRAINING Epoch 41/50 Loss 0.0836 Accuracy 0.5297
TRAINING Epoch 42/50 Loss 0.0824 Accuracy 0.5313
TRAINING Epoch 43/50 Loss 0.0821 Accuracy 0.5260
TRAINING Epoch 44/50 Loss 0.0821 Accuracy 0.5197
TRAINING Epoch 45/50 Loss 0.0807 Accuracy 0.5217
TRAINING Epoch 46/50 Loss 0.0810 Accuracy 0.5230
TRAINING Epoch 47/50 Loss 0.0793 Accuracy 0.5373
TRAINING Epoch 48/50 Loss 0.0787 Accuracy 0.5303
TRAINING Epoch 49/50 Loss 0.0788 Accuracy 0.5367
TRAINING Epoch 50/50 Loss 0.0781 Accuracy 0.5367
Finished Training
```

And test accuracy at 35.44%:

```
test(model, criterion)

Test Loss: 0.0908 Test Accuracy 0.3544
```

## 2. Second try on slight tweaks to those hyperparameters, that is:

NUM_EPOCHS = 50	TRAINING Epoch 40/50 Loss 0.0317 Accuracy 0.8053
	TRAINING Epoch 41/50 Loss 0.0310 Accuracy 0.8093
LEARNING_RATE = 0.001	TRAINING Epoch 42/50 Loss 0.0296 Accuracy 0.8180
	TRAINING Epoch 43/50 Loss 0.0295 Accuracy 0.8213
BATCH_SIZE = 32	TRAINING Epoch 44/50 Loss 0.0284 Accuracy 0.8263
	TRAINING Epoch 45/50 Loss 0.0282 Accuracy 0.8203
RESNET_LAST_ONLY = False	TRAINING Epoch 46/50 Loss 0.0282 Accuracy 0.8247
	TRAINING Epoch 47/50 Loss 0.0265 Accuracy 0.8263
	TRAINING Epoch 48/50 Loss 0.0265 Accuracy 0.8400
I get a train accuracy 84.27%.	TRAINING Epoch 49/50 Loss 0.0257 Accuracy 0.8450
	TRAINING Epoch 50/50 Loss 0.0266 Accuracy 0.8427
	Finished Training

And test accuracy at 56.78%:

```
[27] test(model, criterion)
```

```
Test Loss: 0.0518 Test Accuracy 0.5678
```